2026 9 5

token “ ”
token

1

“ ” [1] token [2]
token Gopher [3] token
RefinedWeb [4] Dolma [5] BaichuanSEED
[6] “ ” token

2

[7] [8]

[9] “ ”
[10] [11]
MiniPile
[12] [13]

3

[14]

“ ” [15]
[16]
[17] [18] “ ”

[19] [20] [21]

4

“ ” “ ” [22] OpenDataVal [23]
[24]
[25] [26] [27] Skill-it
[28] “ ” “ ”
[29] ORCA [30] [31]
[32] LLM [33]

5

DoReMi [34] DoGE [35]
[36] [37] “ ”
SAVA [38] [17]

token token [39] AI
[40] token

6

ChatGLM [41] BaichuanSEED
[6]
CMMLU [42] [43]

“ ” Falcon
RefinedWeb [44] tokenizer

7

tokenizer “ ”

8

token

[45]

9

[1] Kaplan J, McCandlish S, Henighan T, Brown T B, Chess B, Child R, Gray S, Radford A, Wu J, Amodei D. Scaling Laws for Neural Language Models[J/OL]. arXiv (Cornell University), 2020. https://doi.org/10.48550/arxiv.2001.08361.

[2] Hoffmann J, Borgeaud S, Mensch A, Buchatskaya E, Cai T, Rutherford E, de Las Casas D, Hendricks L A, Welbl J, Clark A. Training Compute-Optimal Large Language Models[J/OL]. arXiv (Cornell University), 2022. https://doi.org/10.48550/arxiv.2203.15556.

[3] Rae J W, Borgeaud S, Cai T, Millican K, Hoffmann J, Song F, Aslanides J, Henderson S B, Ring R, Young S. Scaling Language Models: Methods, Analysis & Insights from Training Gopher[J/OL]. arXiv (Cornell University), 2021. https://doi.org/10.48550/arxiv.2112.11446.

[4] Penedo G, Malartic Q, Hesslow D, Cojocaru R, Cappelli A, Alobeidli H, Pannier B, Almazrouei E, Launay J. The RefinedWeb Dataset for Falcon LLM: Outperforming Curated Corpora with Web Data, and Web Data Only[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2306.01116.

[5] Soldaini L, Kinney R, Bhagia A, Schwenk D, Atkinson D, Authur R, Bogin B, Chandu K R, Dumas J, Elazar Y. Dolma: an Open Corpus of Three Trillion Tokens for Language Model Pretraining Research[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.acl-long.840.

[6] Dong G, Pan D, Sun Y, Zhang S, Liang Z, Wu X, Shen Y, Yang F, Sun H, Li T. BaichuanSEED: Sharing the Potential of ExtensivE Data Collection and Deduplication by Introducing a Competitive Large Language Model Baseline[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2408.15079.

[7] Wang Z, Zhong W, Wang Y, Zhu Q, Mi F, Wang B, Shang L, Jiang X, Liu Q. Data Management For Training Large Language Models: A Survey[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2312.01700.

[8] Liu Y, Cao J, Liu C, Ding K, Jin L. Datasets for Large Language Models: A Comprehensive Survey[J/OL]. Research Square, 2024. https://doi.org/10.21203/rs.3.rs-3996137/v1.

[9] Desai M, Pasquetto I V, Jacobs A Z, Card D. An archival perspective on pretraining data[J/OL]. Patterns, 2024. https://doi.org/10.1016/j.patter.2024.100966.

[10] Carlini N, Tramèr F, Wallace E, Jagielski M, Herbert-Voss A, Lee K, Roberts A P, Brown T B, Song D, Úlfar Erlingsson. Extracting Training Data from Large Language Models[J/OL]. arXiv (Cornell University), 2020. https://doi.org/10.48550/arxiv.2012.07805.

[11] Nasr M, Carlini N, Hayase J, Jagielski M, Cooper A F, Ippolito D, Choquette-Choo C A, Wallace E, Tramèr F, Lee K. Scalable Extraction of Training Data from (Production) Language Models[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2311.17035.

[12] Kaddour J. The MiniPile Challenge for Data-Efficient Language Models[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2304.08442.

[13] Elazar Y, Bhagia A, Magnusson I, Ravichander A, Schwenk D, Suhr A, Walsh P, Groeneveld D, Soldaini L, Singh S. What's In My Big Data?[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2310.20707.

[14] Lee K, Ippolito D, Nystrom A, Zhang C, Eck D, Callison-Burch C, Carlini N. Deduplicating Training Data Makes Language Models Better[J/OL]. Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), 2022. https://doi.org/10.18653/v1/2022.acl-long.577.

[15] Vidgen B, Derczynski L. Directions in abusive language training data, a systematic review: Garbage in, garbage out[J/OL]. PLoS ONE, 2020. https://doi.org/10.1371/journal.pone.0243300.

[16] Feng S, Park C Y, Liu Y, Tsvetkov Y. From Pretraining Data to Language Models to Downstream Tasks: Tracking the Trails of Political Biases Leading to Unfair NLP Models[J/OL]. Unknown, 2023. https://doi.org/10.18653/v1/2023.acl-long.656.

[17] Zhang C, Zhong H, Zhang K, Chai C, Wang R, Zhuang X, Bai T, Qiu J, Cao L, Fan J. Harnessing Diversity for Important Data Selection in Pretraining Large Language Models[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2409.16986.

[18] Maini P, Seto S, Bai R, Grangier D, Zhang Y, Jaitly N. Rephrasing the Web: A Recipe for Compute and Data-Efficient Language Modeling[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.acl-long.757.

[19] Tan Z, Li D, Wang S, Beigi A, Jiang B, Bhattacharjee A, Karami M, Li J, Lu C, Liu H. Large Language Models for Data Annotation and Synthesis: A Survey[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.emnlp-main.54.

[20] Long L, Wang R, Xiao R, Zhao J, Ding X, Chen G, Wang H. On LLMs-Driven Synthetic Data Generation, Curation, and Evaluation: A Survey[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.findings-acl.658.

[21] Ding B, Qin C, Zhao R, Luo T, Li X, Chen G, Xia W, Hu J, Luu A T, Joty S. Data Augmentation using LLMs: Data Perspectives, Learning Paradigms and Challenges[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.findings-acl.97.

[22] Hammoudeh Z, Lowd D. Training data influence analysis and estimation: a survey[J/OL]. Machine Learning, 2024. https://doi.org/10.1007/s10994-023-06495-7.

[23] Jiang K F, Liang W, Zou J, Kwon Y. OpenDataVal: a Unified Benchmark for Data Valuation[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2306.10577.

[24] Ki N, Choi H, Chung H W. Data Valuation Without Training of a Model[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2301.00930.

[25] Yeh C, Taly A, Sundararajan M, Liu F, Ravikumar P. First is Better Than Last for Language Data Influence[J/OL]. arXiv (Cornell University), 2022. https://doi.org/10.48550/arxiv.2202.11844.

[26] Thakkar M, Bolukbasi T, Ganapathy S, Vashishth S, Chandar S, Talukdar P. Self-Influence Guided Data Reweighting for Language Model Pre-training[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.18653/v1/2023.emnlp-main.125.

[27] Sow D, Woisetschläger H, Bulusu S, Wang S, Jacobsen H, Liang Y. Dynamic Loss-Based Sample Reweighting for Improved Large Language Model Pretraining[J/OL]. arXiv (Cornell University), 2025. https://doi.org/10.48550/arxiv.2502.06733.

[28] Chen M F, Roberts N, Bhatia K, Wang J, Zhang C, Sala F, Ré C. Skill-it! A Data-Driven Skills Framework for Understanding and Training Language Models[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2307.14430.

[29] Carlini N, Ippolito D, Jagielski M, Lee K, Tramèr F, Zhang C. Quantifying Memorization Across Neural Language Models[J/OL]. arXiv (Cornell University), 2022. https://doi.org/10.48550/arxiv.2202.07646.

[30] Han X, Tsvetkov Y. ORCA: Interpreting Prompted Language Models via Locating Supporting Data Evidence in the Ocean of Pretraining Data[J/OL]. arXiv (Cornell University), 2022. https://doi.org/10.48550/arxiv.2205.12600.

[31] Zhao Y, Du L, Ding X, Xiong K, Sun Z, Jun S, Liu T, Qin B. Deciphering the Impact of Pretraining Data on Large Language Models through Machine Unlearning[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.findings-acl.559.

[32] Chang T A, Rajagopal D, Bolukbasi T, Dixon L, Tenney I. Scalable Influence and Fact Tracing for Large Language Model Pretraining[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2410.17413.

[33] Choe S K, Ahn H, Bae J, Zhao K, Chung Y, Pratapa A, Neiswanger W, Strubell E, Mitamura T, Schneider J. What is Your Data Worth to GPT? LLM-Scale Data Valuation with Influence Functions[J/OL]. Unknown, 2025. https://doi.org/10.52202/085713-4883.

[34] Xie S M, Pham H, Dong X, Du N, Liu H, Lu Y, Liang P, Le Q V, Ma T, Yu A W. DoReMi: Optimizing Data Mixtures Speeds Up Language Model Pretraining[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2305.10429.

[35] Fan S, Pagliardini M, Jaggi M. DoGE: Domain Reweighting with Generalization Estimation[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2310.15393.

[36] Ye J, Liu P, Sun T-C, Zhan J, Zhou Y, Qiu X. Data Mixing Laws: Optimizing Data Mixtures by Predicting Language Modeling Performance[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2403.16952.

[37] Pan R, Zhang D, Zhang H, Pan X, Xu M, Zhang J, Pi R, Wang X, Zhang T. ScaleBiO: Scalable Bilevel Optimization for LLM Data Reweighting[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2406.19976.

[38] Kessler S, Le T-T, Nguyen V. SAVA: Scalable Learning-Agnostic Data Valuation[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2406.01130.

[39] Albalak A, Elazar Y, Xie S M, Longpre S, Lambert N, Wang X, Muennighoff N, Hou B, Pan L, Jeong H. A Survey on Data Selection for Language Models[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2402.16827.

[40] Xu X, Wu Z, Qiao R, Verma A, Shu Y, Wang J, Niu X, He Z, Chen J, Zhou Z. Position Paper: Data-Centric AI in the Age of Large Language Models[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.findings-emnlp.695.

[41] GLM T, :, Zeng A, Xu B, Wang B, Zhang C, Yin D G, Zhang D, Rojas D, Feng G. ChatGLM: A Family of Large Language Models from GLM-130B to GLM-4 All Tools[J/OL]. arXiv (Cornell University), 2024. https://doi.org/10.48550/arxiv.2406.12793.

[42] Li H, Zhang Y, Koto F, Yang Y, Zhao H, Gong Y, Duan N, Baldwin T. CMMLU: Measuring massive multitask language understanding in Chinese[J/OL]. Unknown, 2024. https://doi.org/10.18653/v1/2024.findings-acl.671.

[43] McIntosh T R, Sušnjak T, Arachchilage N, Liu T, Watters P, Halgamuge M N. Inadequacies of Large Language Model Benchmarks in the Era of Generative Artificial Intelligence[J/OL]. arXiv, 2024. https://doi.org/10.48550/arXiv.2402.09880.

[44] Almazrouei E, Alobeidli H, Alshamsi A, Cappelli A, Cojocaru R, Debbah M, Goffinet , Hesslow D, Launay J, Malartic Q. The Falcon Series of Open Language Models[J/OL]. arXiv (Cornell University), 2023. https://doi.org/10.48550/arxiv.2311.16867.

[45] Rao J, Liu X, Yan H, Shen J, Mo H, Dong Y, Yan Z, Wang Z, Lin Z, Meng X. A Data-Centric Perspective on the Lifecycle of Large Language Models[J/OL]. Unknown, 2025. https://doi.org/10.36227/techrxiv.176620610.03288677/v1.